

Model Parameters and Hyper Parameters

```
#initialise parameters and hparameters
hparameters = initialize_hparameters(3, 1, 1)
parameters = initialize_parameters((3,3,3,32), (3,3,32,64), (3,3,64,128), (256, 86528), (10, 256))

"""## **Initialise Parameters function**"""

def initialize_parameters(w1_s, w2_s, w3_s, w4_s, w5_s):

    """ w1_s/w2_s/w3_s is of the form (f,f,n_C_prev,n_C)
        w4_s is of the form (n_hidden_units, length of output from prev layer)
        w5_s is of the form (n_output, length of output from prev layer )

        b1_s/b2_s/b3_s is of the form (1,1,1,n_C)
        b4_s is of the form (n_hidden_units, 1)
        b5_s is of the form (n_output, 1)
    """
    np.random.seed(42)

    W1 = np.random.randn(w1_s[0], w1_s[1], w1_s[2], w1_s[3]) * 0.01
    b1 = np.zeros(shape=(1, 1, 1, w1_s[3]))

    W2 = np.random.randn(w2_s[0], w2_s[1], w2_s[2], w2_s[3]) * 0.01
    b2 = np.zeros(shape=(1,1,1, w2_s[3]))

    W3 = np.random.randn(w3_s[0], w3_s[1], w3_s[2], w3_s[3]) * 0.01
    b3 = np.zeros(shape=(1,1,1, w3_s[3]))

    W4 = np.random.randn(w4_s[0], w4_s[1]) * 0.01
    b4 = np.zeros(shape=(w4_s[0], 1))

    W5 = np.random.randn(w5_s[0], w5_s[1]) * 0.01
    b5 = np.zeros(shape=(w5_s[0], 1))

    parameters = {"W1": W1,
                  "b1": b1,
                  "W2": W2,
                  "b2": b2,
                  "W3": W3,
                  "b3": b3,
                  "W4": W4,
                  "b4": b4,
                  "W5": W5,
                  "b5": b5}

    return parameters

"""## **Initialize hparameter**"""

def initialize_hparameters(f, stride, pad):
    hparameters = {"f": f,
                  "stride": stride,
                  "pad": pad
                  }
    return hparameters
```

Location - helper_functions.py

Location - helper_functions.py

Linear Helper Functions Explained

```
"""## **Unflattening the variable**"""
```

```
def unflatten(A_prev, cache):  
    A = cache  
    a_prev = A_prev.T  
    A_unflat = a_prev.reshape(A.shape)  
    return A_unflat
```

Location – helper_functions.py

```
"""## **One Hot Encoding Function**"""
```

```
def one_hot_encoding(A):  
  
    A = np.squeeze(A.T)  
    a_onehot = np.zeros((A.size , A.max() + 1))  
    a_onehot[np.arange(A.size), A] = 1  
    return a_onehot
```

Location - helper_functions.py

```
"""## **Flatten the input**"""
```

```
def flatten(A_prev):  
    cache = A_prev  
    A_prev_flatten = A_prev.reshape(A_prev.shape[0], -1).T  
    return A_prev_flatten, cache
```

Location – Linear_helper_functions.py

Activation Functions defined using equations

```
def sigmoid(Z):  
    """  
    Implements the sigmoid activation in numpy  
  
    Arguments:  
    Z -- numpy array of any shape  
  
    Returns:  
    A -- output of sigmoid(z), same shape as Z  
    cache -- returns Z as well, useful during backpropagation  
    """
```

```
    A = 1/(1+np.exp(-Z))  
    cache = Z
```

```
    return A, cache
```

```
def relu(Z):  
    """  
    Implement the RELU function.  
  
    Arguments:  
    Z -- Output of the linear layer, of any shape  
  
    Returns:  
    A -- Post-activation parameter, of the same shape as Z  
    cache -- a python dictionary containing "A" ; stored for computing the backward pass efficiently  
    """
```

```
    A = np.maximum(0,Z)  
  
    assert(A.shape == Z.shape)  
  
    cache = Z  
    return A, cache
```

Location – Linear_helper_functions.py

Backward Activation Functions defined for the above activations

```
def sigmoid_backward(dA, cache):  
    """  
    Implement the backward propagation for a single SIGMOID unit.  
  
    Arguments:  
    dA -- post-activation gradient, of any shape  
    cache -- 'Z' where we store for computing backward propagation efficiently  
  
    Returns:  
    dZ -- Gradient of the cost with respect to Z  
    """  
  
    Z = cache  
  
    s = 1/(1+np.exp(-Z))  
    dZ = dA * s * (1-s)  
  
    assert (dZ.shape == Z.shape)  
  
    return dZ
```

Location – Linear_helper_functions.py

For ReLU backwards, rather than implementing the actual derivative for ReLU which basically makes all derivative 0 or 1. I only passed back values that were non zero after implementing relu on the input Z vector.

```
def relu_backward(dA, cache):  
    """  
    Implement the backward propagation for a single RELU unit.  
  
    Arguments:  
    dA -- post-activation gradient, of any shape  
    cache -- 'Z' where we store for computing backward propagation efficiently  
  
    Returns:  
    dZ -- Gradient of the cost with respect to Z  
    """  
  
    Z = cache  
    dZ = np.array(dA, copy=True) # just converting dz to a correct object.  
  
    # When z <= 0, you should set dz to 0 as well.  
    dZ[Z <= 0] = 0  
  
    assert (dZ.shape == Z.shape)  
  
    return dZ
```

Location – Linear_helper_functions.py

Linear forward activation:

```
def linear_forward(A, W, b):
    """
    A -- activations from previous layer (or input data): (size of previous layer, number of examples)
    W -- weights matrix: numpy array of shape (size of current layer, size of previous layer)
    b -- bias vector, numpy array of shape (size of the current layer, 1)

    Z -- the input of the activation function, also called pre-activation parameter
    cache -- a python dictionary containing "A", "W" and "b" ; stored for computing the backward pass efficiently
    """

    Z = W.dot(A) + b

    assert(Z.shape == (W.shape[0], A.shape[1]))
    cache = (A, W, b)

    return Z, cache
```

Location – Linear_helper_functions.py

```
def linear_activation_forward(A_prev, W, b, activation):
    """
    A_prev -- activations from previous layer (or input data): (size of previous layer, number of examples)
    W -- weights matrix: numpy array of shape (size of current layer, size of previous layer)
    b -- bias vector, numpy array of shape (size of the current layer, 1)
    activation -- the activation to be used in this layer, stored as a text string: "sigmoid" or "relu"

    Returns:
    A -- the output of the activation function, also called the post-activation value
    cache -- a python dictionary containing "linear_cache" and "activation_cache";
           stored for computing the backward pass efficiently
    """

    if activation == "sigmoid":
        # Inputs: "A_prev, W, b". Outputs: "A, activation_cache".
        Z, linear_cache = linear_forward(A_prev, W, b)
        A, activation_cache = sigmoid(Z)

    elif activation == "relu":
        # Inputs: "A_prev, W, b". Outputs: "A, activation_cache".
        Z, linear_cache = linear_forward(A_prev, W, b)
        A, activation_cache = relu(Z)

    assert (A.shape == (W.shape[0], A_prev.shape[1]))
    cache = (linear_cache, activation_cache)

    return A, cache
```

Location – Linear_helper_functions.py

Computing Cost:

```
def compute_cost(AL, Y):
    """
    Implement the cost function defined by equation (7).

    Arguments:
    AL -- probability vector corresponding to your label predictions, shape (n_class, number of examples)
    Y -- true "label" vector shape (n_class, number of examples)

    Returns:
    cost -- cross-entropy cost
    """

    m = Y.shape[1]

    # Compute loss from aL and y.
    cost = (1./m) * np.sum(-np.multiply(Y, np.log(AL)).T - np.multiply(1-Y, np.log(1-AL)).T)

    cost = np.squeeze(cost)      # To make sure your cost's shape is what we expect (e.g. this turns [[17]] into 17).

    return cost
```

Location – Linear_helper_functions.py

Linear Backward Activation:

```
def linear_backward(dZ, cache):

    A_prev, W, b = cache
    m = A_prev.shape[1]

    dW = 1./m * np.dot(dZ, A_prev.T)
    db = 1./m * np.sum(dZ, axis = 1, keepdims = True)
    dA_prev = np.dot(W.T, dZ)

    assert (dA_prev.shape == A_prev.shape)
    assert (dW.shape == W.shape)
    assert (db.shape == b.shape)

    return dA_prev, dW, db

def linear_activation_backward(dA, cache, activation):
    """
    Implement the backward propagation for the LINEAR->ACTIVATION layer.

    Arguments:
    dA -- post-activation gradient for current layer l
    cache -- tuple of values (linear_cache, activation_cache) we store for computing backward propagation efficiently
    activation -- the activation to be used in this layer, stored as a text string: "sigmoid" or "relu"

    Returns:
    dA_prev -- Gradient of the cost with respect to the activation (of the previous layer l-1), same shape as A_prev
    dW -- Gradient of the cost with respect to W (current layer l), same shape as W
    db -- Gradient of the cost with respect to b (current layer l), same shape as b
    """
    linear_cache, activation_cache = cache

    if activation == "relu":
        dZ = relu_backward(dA, activation_cache)
        dA_prev, dW, db = linear_backward(dZ, linear_cache)

    elif activation == "sigmoid":
        dZ = sigmoid_backward(dA, activation_cache)
        dA_prev, dW, db = linear_backward(dZ, linear_cache)

    return dA_prev, dW, db
```

Location – Linear_helper_functions.py

Calculating derivative w.r.t linear block

$$J = \frac{1}{m} \sum_{i=1}^m y^{(i)} \log(a_{i2}) + (1-y) \log(1-a_{i2})$$

normalized / avgd.

$$\frac{dJ}{dA_{i2}} \frac{dA_{i2}}{dZ} = dA_{i2} = \frac{-y}{a_{i2}} + \frac{(1-y)}{1-a_{i2}}$$

Diagram illustrating the linear block and its derivatives:

Inputs: dW_5 , db_5 , $dA_{11} = da^{[l-1]}$

Linear block: $a_{i2} = \sigma(W_5 a_{i1} + b_5)$

Output: dA_{i2}

$$dW_5 = \frac{dJ}{dW_5} \times \frac{dW_5}{dA_{i2}} \times dA_{i2}$$

$$= \frac{dA_{i2}}{dW_5} \times \frac{dJ}{dA_{i2}} \rightarrow dA_{i2}$$

$$\underline{dA_{12}} \rightarrow \sigma - \text{relu}$$

$$A_{12} = \sigma(W_5 a_{11} + b_5)$$

$$\rightarrow \frac{dA_{12}}{dW_5} = \underline{A_{11}}$$

$$= \frac{1}{m} \sum_b A_{12} \times (A_{11})^T \rightarrow \text{normalize.}$$

$$\boxed{dW^{[l]} = \frac{1}{m} dz^{[l]} \cdot A^{[l-1]T}}$$

similarly:

$$db = \frac{1}{m} \sum dz^{[l]}$$

$$dA^{[l-1]} = W^{[l]T} dA_{12}^{[l]}$$

→ sending back the derivative

Convolutional Helper Functions Explained

Padding and single step convolution for conv_forward:

```
"""## **Zero Padding Function**"""


def zero_pad(X, pad):
    X_pad = np.pad(X, ((0,0), (pad,pad), (pad,pad), (0,0)), mode = 'constant', constant_values = (0,0))
    return X_pad

"""## **Single Step Convolution**"""

def conv_single_step(a_slice_prev, W, b):
    s = np.multiply(a_slice_prev, W)
    Z = np.sum(s)
    Z = Z + float(b)
    return Z
```


Location – conv_helper_functions.py

<matplotlib.image.AxesImage at 0x7f35d2d35b00>



Convolution Forward:

Logic behind applying mask (image in the right).



```

"""## **Convolution Forward**"""

def conv_forward(A_prev, W, b, hparameters):

    (m, n_H_prev, n_W_prev, n_C_prev) = A_prev.shape

    (f, f, n_C_prev, n_C) = W.shape

    stride = hparameters['stride']
    pad = hparameters['pad']

    n_H = int((n_H_prev - f + 2*pad)/stride) + 1
    n_W = int((n_W_prev - f + 2*pad)/stride) + 1

    Z = np.zeros((m, n_H, n_W, n_C))

    A_prev_pad = zero_pad(A_prev, pad)

    for i in range(m):
        a_prev_pad = A_prev_pad[i,:,:,:]
        for h in range(n_H):
            vert_start = h*stride
            vert_end = h*stride + f

            for w in range(n_W):
                horiz_start = w*stride
                horiz_end = w*stride + f

                for c in range(n_C):
                    a_slice_prev = a_prev_pad[vert_start:vert_end, horiz_start:horiz_end, :]

                    weights = W[:, :, :, c]
                    biases = b[:, :, :, c]
                    Z[i, h, w, c] = conv_single_step(a_slice_prev, weights, biases)
    assert(Z.shape == (m, n_H, n_W, n_C))

    cache = (A_prev, W, b, hparameters)


    return Z, cache

```

Location – conv_helper_functions.py

Pool forward

Max Pool



Max-Pool with a
2 by 2 filter and
stride 2.


```

"""## **Pool Forward Layer**"""
def pool_forward(A_prev, hparameters, mode = "max"):
    """
    Implements the forward pass of the pooling layer

    Arguments:
    A_prev -- Input data, numpy array of shape (m, n_H_prev, n_W_prev, n_C_prev)
    hparameters -- python dictionary containing "f" and "stride"
    mode -- the pooling mode you would like to use, defined as a string ("max" or "average")

    Returns:
    A -- output of the pool layer, a numpy array of shape (m, n_H, n_W, n_C)
    cache -- cache used in the backward pass of the pooling layer, contains the input and hparameters
    """

    # Retrieve dimensions from the input shape
    (m, n_H_prev, n_W_prev, n_C_prev) = A_prev.shape

    # Retrieve hyperparameters from "hparameters"
    f = hparameters["f"]
    stride = hparameters["stride"]

    # Define the dimensions of the output
    n_H = int(1 + (n_H_prev - f) / stride)
    n_W = int(1 + (n_W_prev - f) / stride)
    n_C = n_C_prev

    # Initialize output matrix A
    A = np.zeros((m, n_H, n_W, n_C))

    ### START CODE HERE ###
    for i in range(m):
        # loop over the training examples
        for h in range(n_H):
            # loop on the vertical axis of the output volume
            # Find the vertical start and end of the current "slice" (â‰‰2 lines)
            vert_start = h*stride
            vert_end = h*stride + f

            for w in range(n_W):
                # loop on the horizontal axis of the output volume
                # Find the vertical start and end of the current "slice" (â‰‰2 lines)
                horiz_start = w*stride
                horiz_end = w*stride + f

                for c in range(n_C):
                    # loop over the channels of the output volume

                    # Use the corners to define the current slice on the ith training example of A_prev, channel c. (â‰‰1 line)
                    a_prev_slice = A_prev[i, vert_start:vert_end, horiz_start:horiz_end, c]

                    # Compute the pooling operation on the slice.
                    # Use an if statement to differentiate the modes.
                    # Use np.max and np.mean.
                    if mode == "max":
                        A[i, h, w, c] = np.max(a_prev_slice)
                    elif mode == "average":
                        A[i, h, w, c] = np.mean(a_prev_slice)

    ### END CODE HERE ###

    # Store the input and hparameters in "cache" for pool_backward()
    cache = (A_prev, hparameters)

    # Making sure your output shape is correct
    assert(A.shape == (m, n_H, n_W, n_C))

```

Location – conv_helper_functions.py

Convolution Backwards:

```
"""## **Convolution Backward Pass**"""

def conv_backward(dZ, cache):

    (A_prev, W, b, hparameters) = cache
    (m, n_H_prev, n_W_prev, n_C_prev) = A_prev.shape
    (f, f, n_C_prev, n_C) = W.shape

    stride = hparameters['stride']
    pad = hparameters['pad']

    (m, n_H, n_W, n_C) = dZ.shape

    dA_prev = np.zeros((m, n_H_prev, n_W_prev, n_C_prev))
    dW = np.zeros((f, f, n_C_prev, n_C))
    db = np.zeros((1, 1, 1, n_C))

    A_prev_pad = zero_pad(A_prev, pad)
    dA_prev_pad = zero_pad(dA_prev, pad)

    for i in range(m):
        a_prev_pad = A_prev_pad[i, :, :, :]
        da_prev_pad = dA_prev_pad[i, :, :, :]

        for h in range(n_H):
            for w in range(n_W):
                for c in range(n_C):

                    vert_start = h*stride
                    vert_end = h*stride + f
                    horiz_start = w*stride
                    horiz_end = w*stride + f

                    a_slice = a_prev_pad[vert_start:vert_end, horiz_start:horiz_end, :]

                    da_prev_pad[vert_start:vert_end, horiz_start:horiz_end, :] += W[:, :, :, c]*dZ[i, h, w, c]
                    dW[:, :, :, c] += a_slice*dZ[i, h, w, c]
                    db[:, :, :, c] += dZ[i, h, w, c]

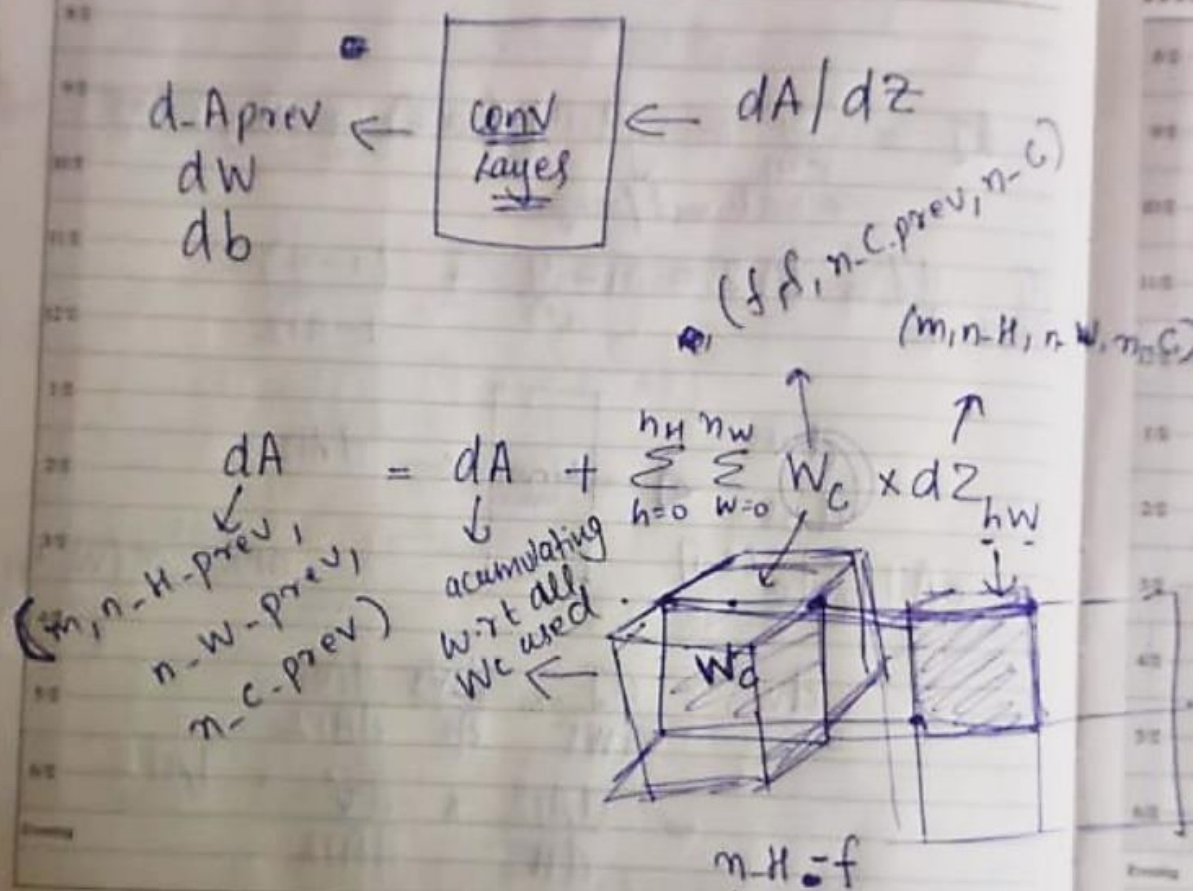
        dA_prev[i, :, :, :] = da_prev_pad[pad:-pad, pad:-pad, :]

    assert(dA_prev.shape == (m, n_H_prev, n_W_prev, n_C_prev))

    return dA_prev, dW, db
```

Location – conv_helper_functions.py

20 November • Monday



Equation for dW_c :

$$dW_c = \sum_{h=0}^{n_H} \sum_{w=0}^{n_W} a_{slice} \times dz_{hw}$$

Annotations for the equation:

- dW_c (pointing to dW_c)
- a_{slice} (pointing to a_{slice})
- dz_{hw} (pointing to dz_{hw})
- n_H (pointing to n_H)
- n_W (pointing to n_W)
- $h=0$ (pointing to $h=0$)
- $w=0$ (pointing to $w=0$)
- a_{slice} (pointing to a_{slice})
- dz_{hw} (pointing to dz_{hw})
- $n-H=f$ (pointing to $n-H=f$)

Diagram showing the accumulation of gradients:

Left side (input to the layer):

- $(m, n-H-1, n-W-1, n-C-1)$
- dA (pointing to $(m, n-H-1, n-W-1, n-C-1)$)

Right side (output of the layer):

- $(f, n-C-1, n-W, n-C)$
- $(m, n-H, n-W, n-C)$

Equation for db :

$$db = \sum_h \sum_w dz_{hw}$$

Annotation for the equation:

- db (pointing to db)

Diagram showing the accumulation of gradients:

Left side (input to the layer):

- $(m, n-H-1, n-W-1, n-C-1)$
- dA (pointing to $(m, n-H-1, n-W-1, n-C-1)$)

Right side (output of the layer):

- $(f, n-C-1, n-W, n-C)$
- $(m, n-H, n-W, n-C)$

similar to Linear back prop.

Pool Backwards:

- 1) Create a mask
- 2) Apply mask to the incoming gradient

```
def create_mask_from_window(x):
    mask = (x == np.max(x))

    return mask

def distribute_value(dZ, shape):
    (n_H, n_W) = shape
    average = np.sum(dZ) / (n_H * n_W)
    a = np.ones((n_H, n_W)) * average

    return a

def pool_backward(dA, cache, mode = "max"):
    """
    Implements the backward pass of the pooling layer

    Arguments:
    dA -- gradient of cost with respect to the output of the pooling layer, same shape as A
    cache -- cache output from the forward pass of the pooling layer, contains the layer's input and hparameters
    mode -- the pooling mode you would like to use, defined as a string ("max" or "average")

    Returns:
    dA_prev -- gradient of cost with respect to the input of the pooling layer, same shape as A_prev
    """

    ### START CODE HERE ###

    # Retrieve information from cache (~1 line)
    (A_prev, hparameters) = cache

    # Retrieve hyperparameters from "hparameters" (~2 lines)
    stride = hparameters["stride"]
    f = hparameters["f"]

    # Retrieve dimensions from A_prev's shape and dA's shape (~2 lines)
    m, n_H_prev, n_W_prev, n_C_prev = A_prev.shape
    m, n_H, n_W, n_C = dA.shape

    # Initialize dA_prev with zeros (~1 line)
    dA_prev = np.zeros(A_prev.shape)

    for i in range(m):
        # loop over the training examples
        # select training example from A_prev (~1 line)
        a_prev = A_prev[i]
        for h in range(n_H):
            # loop on the vertical axis
            for w in range(n_W):
                # loop on the horizontal axis
                for c in range(n_C):
                    # loop over the channels (depth)
                    # Find the corners of the current "slice" (~4 lines)
                    vert_start = h
                    vert_end = vert_start + f
                    horiz_start = w
                    horiz_end = horiz_start + f

                # Compute the backward propagation in both modes.
                if mode == "max":
                    # Use the corners and "c" to define the current slice from a_prev (~1 line)
                    a_prev_slice = a_prev[vert_start:vert_end, horiz_start:horiz_end, c]
                    # Create the mask from a_prev_slice (~1 line)
                    mask = create_mask_from_window(a_prev_slice)
                    # Set dA_prev to be dA_prev + (the mask multiplied by the correct entry of dA) (~1 line)
                    dA_prev[i, vert_start:vert_end, horiz_start:horiz_end, c] += np.multiply(mask, dA[i, h, w, c])

                elif mode == "average":
                    # Get the value a from dA (~1 line)
                    da = dA[i, h, w, c]
                    # Define the shape of the filter as fx (f, f) (~1 line)
                    shape = (f, f)
                    # Distribute it to get the correct slice of dA_prev. i.e. Add the distributed value of da. (~1 line)
                    dA_prev[i, vert_start:vert_end, horiz_start:horiz_end, c] += distribute_value(da, shape)

    ### END CODE ###

    # Making sure your output shape is correct
    assert(dA_prev.shape == A_prev.shape)

    return dA_prev
```

Location – conv_helper_functions.py